

PACKET SNIFFING AND SPOOFING

CS44500 Computer Security

How Packets Are Received

- NIC (Network Interface Card) is a physical or logical link between a machine and a network.
- Each NIC has a MAC address.
- Every NIC on the network will hear all the frames on the wire.
- NIC checks the destination address for every packet, if the address matches the cards MAC address, it is further copied into a buffer in the kernel.

Promiscuous Mode

- The frames that are not destined to a given NIC are usually discarded
- When operating in promiscuous mode, NIC passes every frame received from the network to the kernel. (requires root privileges to be enabled)
- If a sniffer program is registered with the kernel, it will be able to see all the packets
- In Wi-Fi, it is called Monitor Mode

BSD Packet Filter (BPF)

```
struct sock_filter code[] = {
    { 0x28, 0, 0, 0x0000000c }, { 0x15, 0, 8, 0x000086dd },
    { 0x30, 0, 0, 0x00000014 }, { 0x15, 2, 0, 0x00000084 },
    { 0x15, 1, 0, 0x00000006 }, { 0x15, 0, 17, 0x00000011 },
    { 0x28, 0, 0, 0x00000036 }, { 0x15, 14, 0, 0x00000016 },
    { 0x28, 0, 0, 0x00000038 }, { 0x15, 12, 13, 0x00000016 },
    { 0x15, 0, 12, 0x00000800 }, { 0x30, 0, 0, 0x00000017 },
    { 0x15, 2, 0, 0x00000084 }, { 0x15, 1, 0, 0x00000006 },
    { 0x15, 0, 8, 0x00000011 }, { 0x28, 0, 0, 0x00000014 },
    { 0x45, 6, 0, 0x00001fff }, { 0xb1, 0, 0, 0x0000000e },
    { 0x48, 0, 0, 0x0000000e }, { 0x15, 2, 0, 0x00000016 },
    { 0x48, 0, 0, 0x00000010 }, { 0x15, 0, 1, 0x00000016 },
    { 0x06, 0, 0, 0x0000ffff }, { 0x06, 0, 0, 0x00000000 },
};

struct sock_fprog bpf = {
    .len = ARRAY_SIZE(code),
    .filter = code,
};
```


- BPF allows a user-program to attach a filter to the socket, which tells the kernel to discard unwanted packets.
- The filter is often written in human readable format using Boolean operators and is compiled into a pseudo-code and passed to the BPF driver
- An example of the compiled BPF code is shown here.

BSD Packet Filter (BPF)

```
setsockopt(sock, SOL_SOCKET, SO_ATTACH_FILTER, &bpf, sizeof(bpf))
```

- A compiled BPF pseudo-code can be attached to a socket through `setsockopt()`
- When a packet is received by kernel, BPF will be invoked
- An accepted packet is pushed up the protocol stack. See the diagram on the following slide.

Packet Flow With/Without Filters



Packet Sniffing

Packet sniffing describes the process of capturing live data as they flow across a network.

Let's first see how computers receive packets.

Receiving Packets Using Socket

Create the socket

Provide
information about
server

Receive packets

```
// Step ①
int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);

// Step ②
memset((char *) &server, 0, sizeof(server));
server.sin_family = AF_INET;
server.sin_addr.s_addr = htonl(INADDR_ANY);
server.sin_port = htons(9090);

if (bind(sock, (struct sockaddr *) &server, sizeof(server)) < 0)
    error("ERROR on binding");

// Step ③
while (1) {
    bzero(buf, 1500);
    recvfrom(sock, buf, 1500-1, 0,
              (struct sockaddr *) &client, &clientlen);
    printf("%s\n", buf);
}
```


Receiving Packets Using Raw Socket (Packet Sniffing)

Creating a raw socket

Capture all types of packets

```
// Create the raw socket
int sock = socket(AF_PACKET, SOCK_RAW, htons(ETH_P_ALL)); ①

// Turn on the promiscuous mode.
mr.mr_type = PACKET_MR_PROMISC; ②
setsockopt(sock, SOL_PACKET, PACKET_ADD_MEMBERSHIP, &mr, ③
           sizeof(mr));

// Getting captured packets
while (1) {
    int data_size=recvfrom(sock, buffer, PACKET_LEN, 0, ④
                          &saddr, (socklen_t*)sizeof(saddr));
    if(data_size) printf("Got one packet\n");
}
```

Enable the
promiscuous
mode

Wait for packets

Limitation of the Approach

- This program is not portable across different operating systems.
- Setting filters is not easy.
- The program does not explore any optimization to improve performance.
- The PCAP library was thus created.
 - It still uses raw sockets internally, but its API is standard across all platforms. OS specifics are hidden by PCAP's implementation.
 - Allows programmers to specify filtering rules using human readable Boolean expressions.

Packet Sniffing Using the pcap API

```
char filter_exp[] = "ip proto icmp";
```

Filter

```
// Step 1: Open live pcap session on NIC with name eth3  
handle = pcap_open_live("eth3", BUFSIZ, 1, 1000, errbuf); ①  
  
// Step 2: Compile filter_exp into BPF psuedo-code  
pcap_compile(handle, &fp, filter_exp, 0, net); ②  
pcap_setfilter(handle, &fp); ③  
  
// Step 3: Capture packets  
pcap_loop(handle, -1, got_packet, NULL); ④
```

Initialize a raw socket, set the network device into promiscuous mode.

Invoke this function for every captured packet

```
void got_packet(u_char *args, const struct pcap_pkthdr *header,  
               const u_char *packet)  
{  
    printf("Got a packet\n");  
}
```

Processing Captured Packet: Ethernet Header

```
/* Ethernet header */
struct ethheader {
    u_char  ether_dhost[ETHER_ADDR_LEN]; /* destination host address */
    u_char  ether_shost[ETHER_ADDR_LEN]; /* source host address */
    u_short ether_type;                  /* IP? ARP? RARP? etc */
};

void got_packet(u_char *args, const struct pcap_pkthdr *header,
               const u_char *packet)
{
    struct ethheader *eth = (struct ethheader *)packet;
    if (ntohs(eth->ether_type) == 0x0800) { ... } // IP packet
    ...
}
```

The **packet** argument contains a copy of the packet, including the Ethernet header. We typecast it to the Ethernet header structure.

Now we can access the field of the structure

Packet Spoofing

- When some critical information in the packet is forged, we refer to it as packet spoofing.
- Many network attacks rely on packet spoofing.
- Let's see how to send packets without spoofing.

Sending Packets Without Spoofing

```
void main()
{
    struct sockaddr_in dest_info;
    char *data = "UDP message\n";

    // Step 1: Create a network socket
    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);

    // Step 2: Provide information about destination.
    memset((char *) &dest_info, 0, sizeof(dest_info));
    dest_info.sin_family = AF_INET;
    dest_info.sin_addr.s_addr = inet_addr("10.0.2.5");
    dest_info.sin_port = htons(9090);

    // Step 3: Send out the packet.
    sendto(sock, data, strlen(data), 0,
           (struct sockaddr *)&dest_info, sizeof(dest_info));
    close(sock);
}
```

Testing: Use the netcat (nc) command to run a UDP server on 10.0.2.5. We then run the program on the left from another machine. We can see that the message has been delivered to the server machine:



```
seed@Server(10.0.2.5):$ nc -luv 9090
Connection from 10.0.2.6 port 9090 [udp/*] accepted
UDP message
```

Spoofing Packets Using Raw Sockets

There are two major steps in packet spoofing:

- Constructing the packet
- Sending the packet out

Spoofing Packets Using Raw Sockets

```
/******  
Given an IP packet, send it out using a raw socket.  
******/  
void send_raw_ip_packet(struct ipheader* ip)  
{  
    struct sockaddr_in dest_info;  
    int enable = 1;  
  
    // Step 1: Create a raw network socket.  
    int sock = socket(AF_INET, SOCK_RAW, IPPROTO_RAW);  
  
    // Step 2: Set socket option.  
    setsockopt(sock, IPPROTO_IP, IP_HDRINCL,  
               &enable, sizeof(enable));  
  
    // Step 3: Provide needed information about destination.  
    dest_info.sin_family = AF_INET;  
    dest_info.sin_addr = ip->iph_destip;  
  
    // Step 4: Send the packet out.  
    sendto(sock, ip, ntohs(ip->iph_len), 0,  
           (struct sockaddr *)&dest_info, sizeof(dest_info));  
    close(sock);  
}
```

We use *setsockopt()* to enable *IP_HDRINCL* on the socket.

For raw socket programming, since the destination information is already included in the provided IP header, we do not need to fill all the fields

Since the socket type is raw socket, the system will send out the IP packet as is.

Spoofing Packets: Constructing the Packet

Fill in the ICMP Header

```
char buffer[1500];

memset(buffer, 0, 1500);

/*****
  Step 1: Fill in the ICMP header.
  *****/
struct icmpheader *icmp = (struct icmpheader *)
    (buffer + sizeof(struct ipheader));
icmp->icmp_type = 8; //ICMP Type: 8 is request, 0 is reply.

// Calculate the checksum for integrity
icmp->icmp_chksum = 0;
icmp->icmp_chksum = in_cksum((unsigned short *)icmp,
    sizeof(struct icmpheader));
```

Find the starting point
of the ICMP header,
and typecast it to the
ICMP structure

Fill in the ICMP header
fields

Spoofing Packets: Constructing the Packet

Fill in the IP Header

```
/******  
Step 2: Fill in the IP header.  
******/  
struct ipheader *ip = (struct ipheader *) buffer;  
ip->iph_ver = 4;  
ip->iph_ihl = 5;  
ip->iph_ttl = 20;  
ip->iph_sourceip.s_addr = inet_addr("1.2.3.4");  
ip->iph_destip.s_addr = inet_addr("10.0.2.5");  
ip->iph_protocol = IPPROTO_ICMP;  
ip->iph_len = htons(sizeof(struct ipheader) +  
                    sizeof(struct icmphheader));
```

Typecast the buffer to
the IP structure

Fill in the IP header
fields

Finally, send out the packet

```
send_raw_ip_packet (ip);
```

Spoofing UDP Packets

```
memset(buffer, 0, 1500);
struct ipheader *ip = (struct ipheader *) buffer;
struct udphheader *udp = (struct udphheader *) (buffer +
                                                sizeof(struct ipheader));

/*****
  Step 1: Fill in the UDP data field.
  *****/
char *data = buffer + sizeof(struct ipheader) +
              sizeof(struct udphheader);
const char *msg = "Hello Server!\n";
int data_len = strlen(msg);
strncpy (data, msg, data_len);

/*****
  Step 2: Fill in the UDP header.
  *****/
udp->udp_sport = htons(12345);
udp->udp_dport = htons(9090);
udp->udp_ulen = htons(sizeof(struct udphheader) + data_len);
udp->udp_sum = 0; /* Many OSes ignore this field, so we do not
                  calculate it. */
```

← Constructing UDP packets is similar, except that we need to include the payload data now.

Spoofing UDP Packets (continued)

```
/******  
Step 3: Fill in the IP header.  
******/  
  
..... /* Code omitted here; same as that in Listing 12.6 */  
ip->iph_protocol = IPPROTO_UDP; // The value is 17.  
ip->iph_len = htons(sizeof(struct ipheader) +  
                    sizeof(struct udpheader) + data_len);
```

Testing: Use the `nc` command to run a UDP server on `10.0.2.5`. We then spoof a UDP packet from another machine. We can see that the spoofed UDP packet was received by the server machine.

```
seed@Server(10.0.2.5):$ nc -luv 9090  
Connection from 1.2.3.4 port 9090 [udp/*] accepted  
Hello Server!
```

Sniffing and Then Spoofing

- In many situations, we need to capture packets first, and then spoof a response based on the captured packets.
- Procedure (using UDP as example)
 - Use PCAP API to capture the packets of interests
 - Make a copy from the captured packet
 - Replace the UDP data field with a new message and swap the source and destination fields
 - Send out the spoofed reply

Sniffing/Spoofing UDP Packet

```
void spoof_reply(struct ipheader* ip)
{
    const char buffer[1500];
    int ip_header_len = ip->iph_ihl * 4;
    struct udpheader* udp = (struct udpheader *) ((u_char *)ip +
                                                    ip_header_len);

    if (ntohs(udp->udp_dport) != 9999) {
        // Only spoof UDP packet with destination port 9999
        return;
    }

    // Step 1: Make a copy from the original packet
    memset((char*)buffer, 0, 1500);
    memcpy((char*)buffer, ip, ntohs(ip->iph_len));
    struct ipheader * newip = (struct ipheader *) buffer;
    struct udpheader * newudp = (struct udpheader *) (buffer +
                                                         ip_header_len);
    char *data = (char *)newudp + sizeof(struct udpheader);

    // Step 2: Construct the UDP payload, keep track of payload size
    const char *msg = "This is a spoofed reply!\n";
    int data_len = strlen(msg);
    strncpy(data, msg, data_len);
}
```

Sniffing/Spoofing UDP Packet (Continued)

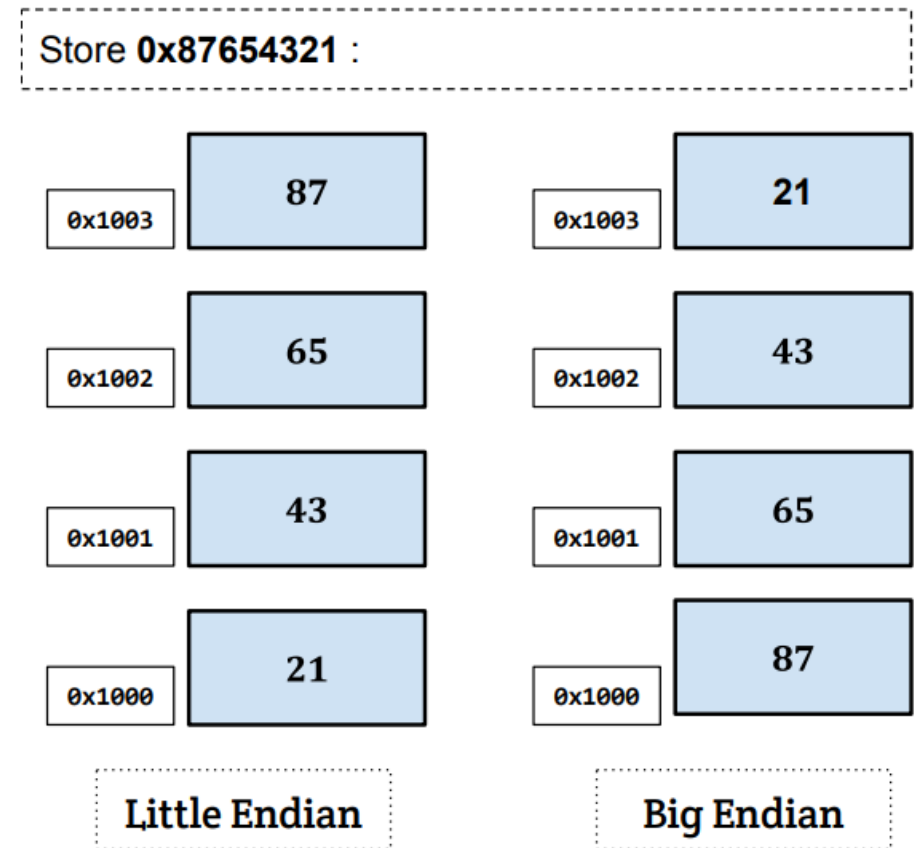
```
// Step 3: Construct the UDP Header
newudp->udp_sport = udp->udp_dport;
newudp->udp_dport = udp->udp_sport;
newudp->udp_ulen = htons(sizeof(struct udphdr) + data_len);
newudp->udp_sum = 0;

// Step 4: Construct the IP header (no change for other fields)
newip->iph_sourceip = ip->iph_destip;
newip->iph_destip = ip->iph_sourceip;
newip->iph_ttl = 50; // Rest the TTL field
newip->iph_len = htons(sizeof(struct iphdr) +
                        sizeof(struct udphdr) + data_len);

// Step 5: Send out the spoofed IP packet
send_raw_ip_packet(newip);
}
```

Endianness

- Endianness: a term that refers to the order in which a given multi-byte data item is stored in memory.
 - **Little Endian**: store the most significant byte of data at the highest address
 - **Big Endian**: store the most significant byte of data at the lowest address



Endianness In Network Communication

- Computers with different byte orders will “misunderstand” each other.
 - Solution: agree upon a common order for communication
 - This is called “network order”, which is the same as big endian order
- All computers need to convert data between “host order” and “network order” .

Macro	Description
<code>htons()</code>	Convert unsigned short integer from host order to network order.
<code>htonl()</code>	Convert unsigned integer from host order to network order.
<code>ntohs()</code>	Convert unsigned short integer from network order to host order.
<code>ntohl()</code>	Convert unsigned integer from network order to host order.

Summary

- Packet sniffing
 - Using raw socket
 - Using PCAP APIs
- Packet spoofing using raw socket
- Sniffing and the spoofing
- Endianness